
A Geometric and Information-Theoretic Interpretation of LayerNorm, FFN, and MHA in Transformers

Anonymous Author(s)

Affiliation

Address

email

Abstract

1 The Transformer architecture has become the foundation of modern artificial intelligence, yet the first-principles reasoning for its design remains surprisingly shallow. 2 Standard explanations—that Layer Normalization (LN) combats “internal covariate shift” or that Multi-Head Attention (MHA) enables parallel processing—are correct but incomplete, failing to explain the deep synergy between the architecture’s 3 components. This paper proposes a new, unified framework that reframes these components through a geometric and information-theoretic lens. We argue that the 4 Transformer block operates on a **symmetrize-and-structure** principle. The process begins with LN acting as a **Geometric Stabilizer**, an isotropic operator that 5 projects token representations onto a fixed, $(d_{\text{model}}-2)$ -dimensional manifold, ensuring dynamic stability. This symmetric “canvas” is then processed by MHA, 6 which we identify as an **Anisotropic Processor** that performs **Axis-Aligned Subspace Decomposition**. A profound implication of this design is that the network is 7 strongly incentivized to **encode meaning in its vector axes**. This is followed by the Feed-Forward Network (FFN), which we frame as a **Manifold-based Information** 8 **Filter** that performs a complexity-reducing deformation to select for relevant features. This unified geometric perspective provides a more powerful, first-principles 9 understanding of the Transformer’s effectiveness and stability. It culminates in a proposal for a fully trainable, anisotropic LayerNorm, which would make the 10 **symmetrize-and-structure** cycle architecturally consistent and biologically plausible, opening a new frontier in model design. The addendum *Summary of* 11 *Theoretical Contributions* provides a consolidated view of our main results. 12 13 14 15 16 17 18 19 20 21 22

23 1 Introduction

24 The Transformer architecture, introduced by Vaswani et al. (2017), has revolutionized natural language 25 processing and beyond. Its success is largely attributed to the self-attention mechanism, which allows 26 for the modeling of global dependencies in a highly parallelizable manner. However, the architecture’s 27 remarkable stability and performance also rely on two other ubiquitous components in every block: a 28 Layer Normalization (LN) layer and a position-wise Feed-Forward Network (FFN).

29 The conventional understanding of these components is largely functional and historical. Layer 30 Normalization (Ba et al., 2016) is typically justified as a technique to stabilize training by normalizing 31 the distribution of activations, thereby reducing “internal covariate shift.” The FFN is explained as 32 being necessary to provide non-linear expressive power, as a stack of linear attention layers would 33 otherwise collapse into a single, less powerful linear transformation.

34 While these explanations are not incorrect, they are unsatisfying. They do not explain, for instance, 35 the precise mathematical synergy between LayerNorm and the $\sqrt{d_k}$ scaling factor in attention. 36 Nor do they provide a compelling rationale for why the specific form of the FFN—an expansion, a




Figure 1: pre-LN Transformer Block

ReLU/GELU activation, and a contraction—is so effective. They describe *what* the components do at a surface level, but not *why* they are designed that way or what their deeper purpose is within the system’s computational logic.

This paper introduces a new perspective that aims to fill these explanatory gaps. We propose a unified framework that interprets LN, MHA, and the FFN not as statistical or algebraic “tricks,” but as fundamental operators performing geometric and information-theoretic transformations. Our core theses are:

1. **Layer Normalization is a Geometric Stabilizer:** Its primary role is to project token representations from the ambient $\mathbb{R}^{d_{\text{model}}}$ space onto a stable, lower-dimensional manifold. The fixed parameters of each LN layer define a unique, time-invariant geometric “canvas” for each processing step, ensuring dynamic stability.
2. **The Feed-Forward Network is an Information Filter:** Its role is to process the representations on this manifold. It does so by first orienting the manifold in a high-dimensional space and then applying a **complexity-reducing deformation** via the activation function, effectively filtering out irrelevant information before summarizing the result.
3. **Multi-Head Attention is an Anisotropic Processor:** Its role is to impose a learned, structured anisotropy on the symmetric space created by LayerNorm. By partitioning the feature space into specialized, parallel subspaces (the heads), it enables a “divide and conquer” strategy that forces the network to encode meaning in the very indices of its vector axes.

This framework provides a more rigorous, first-principles understanding of the Transformer. It explains the synergy between its components, provides a mechanism for its long-context stability, and offers a more powerful vocabulary for analyzing its internal workings.

2 Recap of pre-LN Transformer DecodeBlock

A Transformer DecodeBlock transforms an input sequence representation into an output sequence representation of the same size. While several variants exist, this paper exclusively analyzes the **pre-LN** variant—where Layer Normalization is applied *before* the computational sub-layer—which is common in modern large language models.

- **Input to Block:** $X^{(L-1)} \in \mathbb{R}^{B \times T \times d_{\text{model}}}$. This is a batch of B sequences, each with T token vectors of dimension d_{model} .
- **Output of Block:** $X^{(L)} \in \mathbb{R}^{B \times T \times d_{\text{model}}}$.

The computation within the block proceeds as follows:

1. $X_{\text{norm1}} = \text{LayerNorm1}(X^{(L-1)})$
2. $Y = X^{(L-1)} + \text{MHA}(X_{\text{norm1}})$
3. $Y_{\text{norm2}} = \text{LayerNorm2}(Y)$
4. $X^{(L)} = Y + \text{FFN}(Y_{\text{norm2}})$

73 See *Addendum: DecodeBlock Pseudocode* for a step-by-step exposition.

74 3 The Geometric Stabilizer: A New Role for LayerNorm

75 We begin by re-examining Layer Normalization. In a standard pre-LN Transformer block, every
76 computational sub-layer (both MHA and FFN) is preceded by an LN layer. We argue this is not
77 merely for numerical stability, but to enforce a powerful geometric constraint on the state space.

78 3.1 The Geometry of Normalization: From $\mathbb{R}^{d_{\text{model}}}$ to a $(d_{\text{model}} - 2)$ -Manifold

79 The LayerNorm operation on a single token vector $\mathbf{x} \in \mathbb{R}^{d_{\text{model}}}$ consists of two steps: a normalization
80 and a learned affine transformation. Let us first analyze the normalization step:

$$z = \frac{\mathbf{x} - \mu}{\sqrt{\sigma^2 + \epsilon}}$$

81 where μ and σ^2 are the mean and variance of the components of \mathbf{x} . This operation is not just a
82 statistical rescaling; it is a geometric projection. The output vector \mathbf{z} is not free to point anywhere in
83 $\mathbb{R}^{d_{\text{model}}}$. It is constrained to lie on a specific, lower-dimensional surface defined by two properties:

- 84 1. **Zero Mean:** The sum of its components is zero: $\sum_{i=1}^{d_{\text{model}}} z_i = 0$. This equation defines a
85 hyperplane in $\mathbb{R}^{d_{\text{model}}}$ that passes through the origin. This single linear constraint reduces the
86 degrees of freedom by one.
- 87 2. **Unit Variance / Constant Norm:** The sum of the squares of its components is constant:
88 $\sum_{i=1}^{d_{\text{model}}} z_i^2 = d_{\text{model}} \frac{\sigma^2}{\sigma^2 + \epsilon} \approx d_{\text{model}}$. This equation confines the vector \mathbf{z} to the surface of a
89 hypersphere of radius $\sqrt{d_{\text{model}}}$.

90 The output of the normalization step, \mathbf{z} , must therefore lie on the intersection of a $(d_{\text{model}} - 1)$ -
91 dimensional hyperplane and a $(d_{\text{model}} - 1)$ -dimensional hypersphere. This intersection is itself a
92 sphere of one lower dimension (a $(d_{\text{model}} - 2)$ -sphere), which is a manifold with $d_{\text{model}} - 2$ **intrinsic**
93 **dimensions**.

94 The subsequent affine transformation, $\mathbf{y} = \gamma \odot \mathbf{z} + \beta$, where γ and β are learned parameter vectors,
95 takes this perfect, standardized manifold and performs a layer-specific transformation. The scaling
96 by γ warps the sphere into an ellipsoid, and the shift by β translates it. The final output \mathbf{y} is thus
97 constrained to a specific, learned, $(d_{\text{model}} - 2)$ -dimensional ellipsoid-like manifold.

98 This projection onto a curved, lower-dimensional manifold has a profound computational consequence.
99 While subsequent operations, such as the linear projections in MHA and FFN, are linear with respect
100 to the ambient $\mathbb{R}^{d_{\text{model}}}$ space, they act as powerful **non-linear transformations** on the intrinsic
101 coordinates of the manifold itself. A simple linear shear in a high-dimensional space, for instance,
102 can induce a complex, non-linear warping on the surface of a sphere embedded within it. This
103 interplay between geometric projection and linear algebra is a pervasive and unappreciated source
104 of the Transformer’s non-linear expressive power, operating at every sub-layer of the network. For
105 a concrete 3D illustration of this principle, see the *Addendum: A Geometric Analogy for Implicit*
106 *Non-Linearity*.

107 3.2 Synergy with Attention: Aligning Representational Space with Computation

108 This geometric projection provides a profound explanation for the synergy between LayerNorm and
109 the attention mechanism.

- 110 1. **Factoring out Magnitude:** By forcing all token vectors onto a bounded manifold, LN
111 effectively factors out the vector’s magnitude as a primary information carrier. A vector’s
112 “meaning” is no longer encoded in its length but in its **direction**—its specific location on the
113 manifold.
- 114 2. **Congruence with Attention:** The core operation of attention is the dot product, score
115 $= \mathbf{q} \cdot \mathbf{k}$. We know that $\mathbf{q} \cdot \mathbf{k} = \|\mathbf{q}\| \|\mathbf{k}\| \cos(\text{theta})$. Since the vectors used to
116 generate \mathbf{q} and \mathbf{k} originate from the LN manifold, their norms are constrained. Consequently,
117 the attention score becomes primarily a function of $\cos(\text{theta})$, the cosine similarity.

This reveals a deep design principle: **LayerNorm shapes the latent space to be hyperspherical, and Attention is an operator that naturally measures similarity on a sphere.** The architecture imposes a geometric structure on its data that is perfectly congruent with its primary computational operator. This also provides a first-principles explanation for the $\sqrt{d_k}$ scaling factor. The unit-variance property enforced by LN on the vector components is precisely what causes the dot product’s variance to grow to d_k . The scaling factor is the exact mathematical antidote required to re-normalize the variance of the scores to 1, ensuring stable inputs to the softmax function. This rigorous causal link moves beyond the original explanation in Vaswani et al. (2017), which noted the effect without providing a mechanistic origin in the preceding normalization step.

A Note on Metrics: Geodesic vs. Euclidean: This alignment between geometry and computation becomes even more profound when we consider the concept of distance. In the high-dimensional Euclidean space, distance is a noisy and often uninformative metric. However, by projecting vectors onto a unit hypersphere, the model creates a space where the most meaningful measure of separation is the **geodesic distance**—the shortest path between two points along the curved surface of the sphere.

For two unit vectors on a hypersphere, this geodesic distance is simply the angle, θ , between them. This angle is directly and linearly recoverable from the dot product that attention computes: $\theta = \arccos(\mathbf{q} \cdot \mathbf{k})$. This is a crucial insight. It means the model has discovered a geometry where its fundamental algebraic operation (the dot product) is a direct, uncorrupted measure of true geometric distance. The “curse of dimensionality” is sidestepped by adopting a geometry where angular separation *is* distance, perfectly aligning the representational space with the computational mechanism.

3.3 The Fixed Manifold: A Mechanism for Dynamic Stability

The geometric view of Layer Normalization provides its most significant explanatory power when we consider the dynamics of autoregressive generation. A key challenge for any generative model is maintaining coherence and stability over long sequences. The “Geometric Stabilizer” framework reveals a powerful mechanism by which the Transformer achieves this, though the precise nature of the mechanism depends on the model’s operational context.

The core principle is that the gamma (γ) and beta (β) parameter vectors for each LN layer are learned during training and are **fixed** during inference. Crucially, these parameters are shared across all token positions within a layer, meaning a single, layer-specific transformation is applied to every token in a sequence. This ensures that the geometric manifold each LN layer projects onto is also **fixed and constant across all time steps**. This fixed manifold acts as a powerful regularizer, but its role is best understood by distinguishing between two scenarios: the idealized, infinite-context Transformer and the practical, sliding-window Transformer.

Stability in the Idealized Transformer: Process Consistency

In a standard decoder-only Transformer with an infinite context window, the causal attention mask ensures that the representation of a token at position p is a function only of the tokens at positions $1 \dots p$. Consequently, the representations of previously generated tokens are immutable; they are computed once and do not change as new tokens are appended. This immutability is the foundation of the KV Cache optimization used in modern inference engines.

In this idealized setting, the stability provided by the fixed manifold is not about correcting the “drift” of a single token’s representation over time. Instead, it ensures **process consistency**. At any given time step t , the model processes the entire sequence of $t-1$ tokens in parallel. The fixed manifold at layer l , M_l , guarantees that the vectors for $\text{tok}_1, \text{tok}_2, \dots$, and $\text{tok}_{\{t-1\}}$ are all projected onto the same geometric surface before being processed by the MHA or FFN. This creates a stable and predictable computational regime, ensuring the “rules of the game” for processing any token, at any position, remain constant. Our framework provides a novel, geometric explanation for this architectural stability, moving beyond purely statistical or empirical observations.

Stability in the Practical Transformer: Taming Representational Drift

In practice, to handle sequences longer than their training context length (e.g., 4096 tokens), models are often deployed with a **sliding context window**. In this scenario, as new tokens are generated, the oldest tokens are evicted from the context. This means the causal history for every token within the window changes at every step, forcing a re-computation of their representations.

Here, **representational drift is a real and significant phenomenon**. The vector for a token that is at position p in the window at time t will be different from its representation at time $t+1$ when it has shifted to position $p-1$. In this challenging but common scenario, the role of the fixed manifold as a **geometric stabilizer** becomes even more critical. The repeated re-projection of these constantly changing token vectors onto the same fixed manifold at each layer and each time step acts as a powerful regularizer. It prevents the model’s internal state from spiraling into chaotic or exploding regions of the embedding space, thereby ensuring coherence over extremely long generations.

In both the idealized and practical settings, the fixed manifold of Layer Normalization is a cornerstone of the Transformer’s dynamic stability, providing a robust geometric foundation for coherent, long-form generation. This two-part explanation, distinguishing between the idealized and practical cases, offers a more nuanced understanding of stability than is typically discussed, grounding it in the specific geometric constraints of the architecture.

4 The Information Filter: A New Role for the FFN

With a stable geometric canvas provided by LayerNorm, we can now reinterpret the role of the Feed-Forward Network (FFN). We argue the FFN is not a generic function approximator but a specialized, **manifold-based information filter**, whose power derives from the Universal Approximation Theorem, which states that a two-layer network with a non-linear activation can approximate any continuous function, given sufficient width.

4.1 The FFN as a Manifold Processor

The FFN sub-layer computes $\text{FFN}(y)$, where y is a point on the input manifold M_{in} . The FFN itself consists of two linear layers separated by a non-linearity (e.g., GELU).

1. **Expansion** ($h = \text{GELU}(y * W_1 + b_1)$): The first linear layer, W_1 , expands the dimensionality from d_{model} to d_{ffn} (typically $4 * d_{\text{model}}$). Geometrically, this is a learned **orientation module**. It takes the input manifold M_{in} and performs a linear transformation (rotation, scaling, translation) to place it in a specific orientation within the higher-dimensional d_{ffn} space. This new manifold is M_h .
2. **Contraction** ($o = h * W_2 + b_2$): The second linear layer, W_2 , projects the result back down to d_{model} . Geometrically, this is a **summarization module** that learns to interpret the result of the intermediate processing step.

The critical question is: what is the purpose of the non-linear activation between these two steps?

4.2 The Activation as a Complexity-Reducing Deformation

The conventional answer, “to provide non-linearity,” is true but insufficient. It does not explain why a simple, information-destroying function like ReLU ($\max(0, x)$) or GELU (which approximates ReLU) is so effective.

We propose an information-theoretic interpretation: the activation function acts as a **fixed geometric filter** that performs a **complexity-reducing deformation** on the manifold M_h .

- **Mechanism:** Functions like ReLU and GELU are non-linear, but they have a specific character: they map a large portion of their input domain (e.g., all negative numbers) to a single, simple output (zero). This is an act of **lossy compression**.
- **Geometric Interpretation:** The activation function acts as a fixed stencil on the d_{ffn} space. The W_1 orientation module learns to position the manifold M_h such that the coordinates corresponding to **irrelevant or noisy features** for a given token fall into the filter’s “discard” region (i.e., the negative half-spaces). The activation function then geometrically collapses these regions of the manifold towards the origin, effectively erasing that information.
- **Information-Theoretic Consequence:** This deformation is a direct, mechanistic implementation of **Kolmogorov complexity reduction**. A vector with many zeros is algorithmically simpler to describe than a dense vector. The FFN learns to simplify its representation by strategically destroying information.

220 4.3 The FFN as a Learned, Context-Dependent Feature Selector

221 This synthesis provides a powerful, purposeful model for the FFN’s function. The FFN is a two-stage
222 feature selection mechanism:

- 223 1. The **learned orientation module** (W_1) intelligently positions the token’s representation
224 relative to the filter.
- 225 2. The **fixed geometric filter (the activation)** performs a complexity-reducing deformation,
226 discarding the features that W_1 has positioned in the “off” region.

227 The FFN’s role is to compute a filtered update vector. The residual connection then adds this
228 intelligently simplified update to the original token representation. This allows the model to refine a
229 token’s meaning based only on the essential features that survived the filtering process.

230 5 The Anisotropic Processor: A Geometric View of Multi-Head Attention

231 Having established Layer Normalization as a “Geometric Stabilizer” and the FFN as a “Manifold-
232 based Information Filter,” we now turn our attention to the Multi-Head Attention (MHA) sub-block.
233 The conventional view of MHA focuses on its ability to weigh the importance of different tokens. We
234 propose a complementary geometric interpretation that frames MHA as a structured, **anisotropic**
235 **processor**. This view reveals a deep design principle in the Transformer block: a repeating rhythm
236 of isotropic regularization followed by anisotropic specialization, a **symmetrize-and-structure**
237 cycle that appears fundamental to its ability to process information.

238 5.1 The Isotropic-Anisotropic Dichotomy: A Symmetrize-and-Structure Principle

239 The Transformer block exhibits a remarkable architectural pattern. Every major computational unit
240 (MHA, FFN) is preceded by an operation (LayerNorm) with a fundamentally opposite geometric
241 character.

- 242 1. **Symmetrize (Isotropic Projection):** As established in Section 2, LayerNorm is an **isotropic**
243 operator. It treats all d_{model} dimensions of a token vector identically, projecting the vector
244 onto a symmetric, hyperspherical manifold. This act of regularization creates a clean,
245 standardized, and symmetric “canvas” by factoring out noisy variations in magnitude and
246 placing all representations on an equal geometric footing.
- 247 2. **Structure (Anisotropic Processing):** The MHA block, which acts upon this clean canvas,
248 is a fundamentally **anisotropic** operator. As we will detail, its core mechanism intentionally
249 breaks the symmetry of the d_{model} space, imposing a specific, learned structure to enable
250 specialized computation.

251 This **symmetrize-and-structure** sequence is a powerful design principle. The isotropic projection
252 provides stability and a common geometric ground, while the subsequent anisotropic processing
253 allows for a sophisticated, “divide and conquer” approach to computation.

254 5.2 The Geometry of Anisotropy: Axis-Aligned Subspace Decomposition

255 The anisotropy of MHA is not random; it is a highly structured feature that originates from the
256 **reshape** operation used to create the multiple heads. Before the attention calculation, the d_{model} -
257 dimensional representation of each token is reshaped into H separate head vectors, each of dimension
258 d_v (where $d_{\text{model}} = H \times d_v$). This is a hard, **axis-aligned partitioning** of the vector space. For
259 example, axes 0 to d_v-1 are grouped into Head 1, axes d_v to $2d_v-1$ into Head 2, and so on.

260 A profound implication of this design is that the network is **strongly incentivized to learn to encode**
261 **meaning in its axis indices**. For the partitioning to be useful, the optimization process must discover
262 a consistent “data layout,” placing semantically related features onto contiguous axes that will be
263 processed together by a dedicated attention head.

264 This architectural choice provides an elegant solution to a core challenge of operating in high-
265 dimensional spaces. In a space with thousands of dimensions, two randomly chosen vectors are
266 almost always nearly orthogonal. This “curse of orthogonality” would make it difficult for a single,

global dot-product attention mechanism to learn meaningful similarity scores. The MHA’s anisotropic partitioning circumvents this problem. By decomposing the high-dimensional space into a set of lower-dimensional subspaces (the heads), it creates “pockets” of non-trivial geometry where dense, meaningful, non-orthogonal relationships can be learned effectively. The anisotropy can therefore be seen as a **highly effective strategy** for imposing a learnable, interactive structure on an otherwise uniform and non-interactive high-dimensional space.

5.3 Attention as Relational Computation in Specialized Bases

Within each of these anisotropically defined subspaces, the MHA block performs another layer of specialization. The projection of the input vectors into Query (Q), Key (K), and Value (V) matrices can be interpreted as a learned **change of basis**.

Each head does not compute attention in the general-purpose basis of the main residual stream. Instead, it learns three specialized linear projections (W^Q, W^K, W^V) that transform the input into a new coordinate system optimized for its specific task. This allows different heads to specialize in detecting different kinds of inter-token relationships:

- **Head 1** might learn a basis optimized for resolving syntactic dependencies, where the dot product between Q and K vectors effectively measures grammatical agreement.
- **Head 2** might learn a basis optimized for tracking semantic similarity or opposition.
- **Head 3** might learn a basis optimized for anaphora resolution, where the Q vector for a pronoun like “it” is projected to be highly similar to the K vector of its antecedent.

This view frames the MHA block as performing a two-level decomposition: first, an anisotropic partitioning of the feature space into subspaces, and second, a further projection within each subspace into specialized computational bases. This allows the model to analyze the input sequence from multiple, parallel, and specialized relational perspectives simultaneously. See *Addendum: The Mechanics of Head Specialization* for a step-by-step exposition.

5.4 Future Research Directions

This geometric interpretation of MHA as an anisotropic processor opens up several new and promising avenues for research.

Probing for Axis-Encoded Meaning: The central hypothesis that the network learns to place semantically related features on contiguous axes is empirically testable. One could analyze the activation patterns of different heads when processing specific linguistic phenomena. A more direct test would involve targeted ablation studies: measure the performance degradation when shuffling axes *within* a single head’s designated block versus shuffling the same number of axes *across* different head-blocks. Our framework would predict a significantly larger performance drop in the latter case, which would provide strong evidence for learned, axis-dependent feature grouping.

Generalizing the Anisotropy with “Soft” Heads: The current reshape operation creates a hard, block-diagonal partitioning. We can model this partitioning function as a periodic “impulse train,” where an axis belongs 100% to one head and 0% to all others. This framing invites a natural generalization: what if we used other periodic functions to define the heads? For example, using overlapping sine or triangular waves could create “soft” heads, where an axis might contribute 80% of its activation to Head h and 20% to Head $h+1$. This would allow for more flexible and potentially more powerful feature grouping, moving from a discrete to a continuous model of subspace decomposition.

Connecting Head Specialization to Manifold Geometry: An open theoretical question is how the anisotropic processing of MHA interacts with the geometric manifold created by the preceding LayerNorm. Does the MHA block as a whole learn to operate on the global geometry of the manifold? Or do individual heads learn to specialize on different regions or curvatures of the input manifold, with the axis-aligned partitioning serving as a routing mechanism to direct different parts of the manifold to the appropriate specialist head? Answering this would provide a truly unified geometric theory of the entire Transformer block.

5.5 The Manifold and the Grid: A Productive Incommensurability

Our framework reveals a final, subtle tension at the heart of the `symmetrize-and-structure` cycle. The `symmetrize` step (LayerNorm) projects token representations onto a smooth, $(d_{\text{model}}-2)$ -dimensional manifold. The `structure` step (MHA) then processes this manifold using a rigid, d_{model} -dimensional grid, partitioned into H discrete heads. There is a fundamental **incommensurability** between the dimensionality of the data and the dimensionality of the processor.

We hypothesize that this mismatch is not a design flaw but a powerful form of implicit regularization. It is mathematically impossible to perfectly and evenly partition a $(d_{\text{model}}-2)$ -dimensional object into subspaces defined on a d_{model} -dimensional grid. This geometric friction prevents the model from learning a brittle, “too-perfect” alignment between its learned features and the arbitrary boundaries of the attention heads. Instead, the model is forced to learn more **distributed and robust representations** that are resilient to the slight geometric aliasing that occurs at the LN/MHA interface. This productive tension between the continuous geometry of the manifold and the discrete architecture of the grid may be a key, unappreciated reason for the robustness and generalizability of the Transformer architecture.

6 Discussion and Future Work

A Unified View: The Symmetrize-and-Structure Principle: The geometric and information-theoretic framework presented here unifies the roles of LayerNorm, Multi-Head Attention, and the FFN into a single, coherent narrative. The Transformer block is not an arbitrary collection of layers but a sophisticated computational engine operating on a `symmetrize-and-structure` principle. The process begins with an isotropic projection by LayerNorm, which acts as a **Geometric Stabilizer** to create a stable, symmetric manifold. This is followed by the **Anisotropic Processor** of MHA, which imposes a learned, axis-aligned structure on this manifold to enable specialized, parallel computation. Finally, after another stabilization step, the **Information Filter** of the FFN performs a complexity-reducing deformation to perform context-dependent feature selection.

Consilience with the Renormalization Group: This perspective enriches other theoretical models, such as the view of deep networks as performing a Renormalization Group (RG) flow. In that analogy, the layers correspond to a change in scale. Our framework provides the concrete mechanism for this flow: the `symmetrize-and-structure` cycle is a robust implementation of an RG step. The isotropic projections (LayerNorm) act as the coarse-graining or regularization operators that stabilize the flow, while the anisotropic processors (MHA, FFN) act as the update-computation engines that calculate the new effective couplings for the next scale. This suggests the Transformer has convergently discovered a fundamental principle of information processing in complex, multi-scale systems.

A Generative Framework for Future Work: This framework also opens several new avenues for empirical research and model design. Because the core geometric claims of our framework are derived mathematically from the architecture’s definitions, the primary value of empirical testing is not to “validate” these truths, but to use them as a lens to probe the model’s learned behavior. Such experiments are likely to yield serendipitous inspirations and reveal further puzzles about the nature of the learned representations.

Probing the Latent Geometry: As detailed in Section 4.4, the geometric view of MHA inspires concrete proposals for probing axis-encoded meaning and designing novel “soft head” architectures. Beyond MHA-specific inquiries, our framework’s most direct global prediction is that Layer Normalization reduces the effective dimensionality of the latent space by two. This can be empirically verified by computing the covariance matrix of post-LN activation vectors; our theory predicts this matrix will have **two** eigenvalues that are orders of magnitude smaller than the others, corresponding to the zero-mean and constant-norm constraints.

Designing New Components: A deeper geometric understanding can also inspire the design of entirely new normalization or activation layers based on explicit geometric principles. Furthermore, this framework provides a new vocabulary for failure analysis, allowing us to ask if model errors can be understood as instances where representations fall “off” the learned manifolds or are improperly processed by the anisotropic components.

A Final, Speculative Projection: The Evolution of Structure Itself: Our framework’s re-interpretation of MHA’s fixed partitioning as a specific “impulse train”-like function naturally invites a generalization to “soft heads,” where this rigid structure is replaced by a learnable, parameterized function. This step alone, moving a core architectural heuristic from a static rule to a dynamic parameter, points toward a more profound future. It is conceivable that in a continuous, AI-supervised training regime, the partitioning heuristic itself could become an object of evolution, with a “supervisor” AI dynamically modifying the “worker” AI’s internal structure to best suit the data domain it currently faces. This vision culminates in the proposal for a fully trainable, anisotropic LayerNorm, as detailed in the $\omega + 1$ addendum. Such a component, which would learn the `symmetrize` operation through a parallel, iterative process, would resolve the final architectural tension of the Transformer block, making the `symmetrize-and-structure` cycle fully learnable, architecturally consistent, and more biologically plausible. This points towards a future where our models do not just learn *within* a fixed architecture, but learn *how to architect themselves* for optimal information processing.

7 Conclusion

We have proposed a new, unified framework for understanding the Layer Normalization, Multi-Head Attention, and Feed-Forward Network components of the Transformer architecture. We argue that their roles are not merely statistical or algebraic, but are fundamentally geometric and information-theoretic, operating in a `symmetrize-and-structure` cycle.

LayerNorm acts as a **Geometric Stabilizer**, performing an isotropic projection of representations onto fixed, layer-specific manifolds to ensure dynamic stability. Multi-Head Attention then acts as an **Anisotropic Processor**, imposing a learned, axis-aligned structure on this symmetric space to enable specialized, parallel computation. Finally, the FFN acts as a **Manifold-based Information Filter**, using a complexity-reducing deformation to perform context-dependent feature selection.

This unified perspective provides a more satisfying, first-principles explanation for the Transformer’s design and effectiveness. It reveals a deep architectural logic that balances symmetry with structure, and stability with specialization. This framework offers a powerful new lens for future research into the mechanisms of deep learning, suggesting new ways to probe, analyze, and design the next generation of intelligent systems.

References

- Ba, J. L., Kiros, J. R., & Hinton, G. E. (2016). Layer normalization. *arXiv preprint arXiv:1607.06450*.
- He, K., Zhang, X., Ren, S., & Sun, J. (2016). Deep residual learning for image recognition. In *Proceedings of the IEEE conference on computer vision and pattern recognition* (pp. 770-778).
- Hendrycks, D., & Gimpel, K. (2016). Gaussian Error Linear Units (GELUs). *arXiv preprint arXiv:1606.08415*.
- Mehta, P., & Schwab, D. J. (2014). An exact mapping between the Variational Renormalization Group and Deep Learning. *arXiv preprint arXiv:1410.3831*.
- Nair, V., & Hinton, G. E. (2010). Rectified linear units improve restricted boltzmann machines. In *Proceedings of the 27th international conference on machine learning (ICML-10)* (pp. 807-814).
- Vaswani, A., Shazeer, N., Parmar, N., Uszkoreit, J., Jones, L., Gomez, A. N., ... & Polosukhin, I. (2017). Attention is all you need. *Advances in neural information processing systems*, 30.
- Xiong, R., Yang, Y., He, D., Zheng, S., Zheng, S., Xing, C., ... & Liu, T. Y. (2020). On Layer Normalization in the Transformer Architecture. In *Proceedings of the 37th International Conference on Machine Learning, PMLR 119* (pp. 10525-10535).

$\omega + 1$: And One More Thing: a Trainable LayerNorm?

Our analysis has revealed the Transformer block to be a sophisticated engine operating on a `symmetrize-and-structure` principle. We have identified Layer Normalization as the **Geometric Stabilizer**, an isotropic operator that projects representations onto a fixed manifold, thereby creating

414 a symmetric canvas for computation. This is followed by the **Anisotropic Processor** of Multi-Head
415 Attention, which imposes a learned, structured partitioning on this canvas. This isotropic ->
416 anisotropic cycle is a deep and powerful design pattern.

417 Yet, it presents a final, subtle tension. The *symmetrize* step is performed by a fixed, analytical, and
418 global algorithm, making it a rigid outlier in an otherwise learnable, parallel, and structured system.
419 This invites a final, speculative question: Now that we understand the *function* of Layer Normalization,
420 can we design a better *implementation*—one that is more harmonious with the architecture it serves
421 and more congruent with the principles of neural computation?

422 We propose that the next frontier is to replace the standard LayerNorm algorithm with a learnable,
423 parallel, and iterative block. This “symmetrizing block” would be a neural network module in its own
424 right, designed to perform the same geometric function of projecting a vector onto a stable manifold,
425 but through a different mechanism. Crucially, this block would be designed to be **anisotropic**
426 **from the start**. It would be composed of H parallel heads, each responsible for a subset of the
427 d_{model} dimensions, mirroring the structure of the MHA block it precedes. This would resolve the
428 architectural tension, transforming the isotropic -> anisotropic cycle into a more elegant and
429 consistent anisotropic *symmetrize* -> anisotropic structure rhythm.

430 Such a proposal is motivated not only by the pursuit of architectural elegance but also by biological
431 plausibility. It is difficult to imagine a biological system implementing a global “compute mean and
432 variance” operation. A parallel, iterative algorithm, where local computational units communicate
433 to converge on a global property, is a far more likely model for how a brain might achieve a similar
434 normalization function.

435 The engineering challenges would be significant. A randomly initialized block would likely re-
436 introduce the very instabilities that Pre-LN architectures solve, necessitating a careful pre-training
437 regimen (perhaps by first learning to mimic the standard LN function) and a slow learning rate
438 schedule. However, the parameters of this block could be learned once and shared universally, not
439 just across all layers of a single model, but potentially as a universal, pre-trained “symmetrizing chip”
440 for all Transformers.

441 This is not a finished idea, but a research direction opened up by the geometric framework. It is an
442 invitation to move beyond analyzing the architectures we have and begin designing new components
443 based on the principles we have uncovered. It is a step towards a future where the fundamental
444 building blocks of our models are not just effective, but are also learnable, principled, and perhaps, a
445 little closer to the way intelligence is organized in the natural world.

446 **Addendum: DecodeBlock Pseudocode**

447 This addendum details the step-by-step computations within a single DecodeBlock of a decoder-only
448 Transformer. We use TeX notation and explicitly state the dimensions (rows \times columns) of all data
449 and weight matrices.

450 **Notation and Dimensions**

451 Let the primary hyperparameters of the model be:

- 452 • B : Batch size (number of sequences processed in parallel).
- 453 • T : Sequence length (number of tokens in each sequence).
- 454 • d_{model} : The main embedding dimension of the model.
- 455 • H : The number of parallel attention heads.
- 456 • d_k : The dimension of the Query and Key vectors for each head.
- 457 • d_v : The dimension of the Value vector for each head.

458 For standard Transformer architectures, the dimensions are constrained such that the total capacity
459 across all heads equals the model’s main dimension: $H \cdot d_k = H \cdot d_v = d_{\text{model}}$.

460 The DecodeBlock

461 A DecodeBlock transforms an input sequence representation into an output sequence representation
462 of the same size, following the pre-normalization variant common in modern GPT-style models.

- 463 • **Input to Block:** $X^{(L-1)} \in \mathbb{R}^{B \times T \times d_{\text{model}}}$. This is a batch of B sequences, each with T token
464 vectors of dimension d_{model} .
- 465 • **Output of Block:** $X^{(L)} \in \mathbb{R}^{B \times T \times d_{\text{model}}}$.

466 The computation within the block proceeds as follows:

- 467 1. $X_{\text{norm1}} = \text{LayerNorm1}(X^{(L-1)})$
- 468 2. $Y = X^{(L-1)} + \text{MHA}(X_{\text{norm1}})$
- 469 3. $Y_{\text{norm2}} = \text{LayerNorm2}(Y)$
- 470 4. $X^{(L)} = Y + \text{FFN}(Y_{\text{norm2}})$

471 We now detail the computations within the Multi-Head Attention (MHA) block.

472 Detailed Computation within the Multi-Head Attention (MHA) Block

473 The MHA block takes the normalized input X_{norm1} and produces an output of the same dimension.

- 474 • **Input to MHA:** $X_{\text{in}} \in \mathbb{R}^{B \times T \times d_{\text{model}}}$ (where $X_{\text{in}} = X_{\text{norm1}}$).
- 475 • **Output of MHA:** $\text{MHA}_{\text{out}} \in \mathbb{R}^{B \times T \times d_{\text{model}}}$.

476 MHA Sub-block 2: Learned Projections for Q, K, and V

477 The first step is to project the input X_{in} into three separate matrices: Query (Q), Key (K), and Value
478 (V). This is done using three distinct, learned linear projection weight matrices.

479 • Weight Matrices:

- 480 – $W^Q \in \mathbb{R}^{d_{\text{model}} \times d_{\text{model}}}$ (rows \times columns: $d_{\text{model}} \times d_{\text{model}}$)
- 481 – $W^K \in \mathbb{R}^{d_{\text{model}} \times d_{\text{model}}}$ (rows \times columns: $d_{\text{model}} \times d_{\text{model}}$)
- 482 – $W^V \in \mathbb{R}^{d_{\text{model}} \times d_{\text{model}}}$ (rows \times columns: $d_{\text{model}} \times d_{\text{model}}$)

483 • Computation:

$$Q' = X_{\text{in}} W^Q \quad | \quad K' = X_{\text{in}} W^K \quad | \quad V' = X_{\text{in}} W^V$$

- 484 • **Resulting Dimensions:** The resulting matrices Q' , K' , and V' each have the dimension
485 $\mathbb{R}^{B \times T \times d_{\text{model}}}$.

- 486 • **Prepare for Multi-Head Processing:** These matrices are then reshaped and transposed to
487 split the d_{model} dimension across the H attention heads.

- 488 – $Q = \text{reshape}(Q') \rightarrow \mathbb{R}^{B \times T \times H \times d_k} \xrightarrow{\text{transpose}} \mathbb{R}^{B \times H \times T \times d_k}$
- 489 – $K = \text{reshape}(K') \rightarrow \mathbb{R}^{B \times T \times H \times d_k} \xrightarrow{\text{transpose}} \mathbb{R}^{B \times H \times T \times d_k}$
- 490 – $V = \text{reshape}(V') \rightarrow \mathbb{R}^{B \times T \times H \times d_v} \xrightarrow{\text{transpose}} \mathbb{R}^{B \times H \times T \times d_v}$

491 MHA Sub-block 1: Scaled Dot-Product Attention

492 This core operation is performed independently for each head in parallel.

- 493 • **Inputs:** $Q, K \in \mathbb{R}^{B \times H \times T \times d_k}$ and $V \in \mathbb{R}^{B \times H \times T \times d_v}$.
- 494 • **Step 1: Compute Attention Scores:** The dot product between every query vector and every
495 key vector is computed.

$$\text{Scores} = Q \cdot K^T$$

496 The matrix multiplication is performed on the last two dimensions. The dimensions are:
497 $(\mathbb{R}^{B \times H \times T \times d_k}) \cdot (\mathbb{R}^{B \times H \times d_k \times T}) \rightarrow \mathbb{R}^{B \times H \times T \times T}$. The resulting Scores matrix contains,
498 for each head, a $T \times T$ matrix of similarity scores between all pairs of tokens in the sequence.

- 499 • **Step 2: Scale, Mask, and Softmax:** The scores are scaled to stabilize gradients, masked
500 to enforce causality (i.e., a token cannot attend to future tokens), and normalized into a
501 probability distribution using softmax.

$$\text{AttentionWeights} = \text{softmax} \left(\frac{\text{Scores}}{\sqrt{d_k}} + M \right)$$

- 502 – The mask $M \in \mathbb{R}^{T \times T}$ is a matrix where $M_{ij} = -\infty$ for $j > i$ (upper triangle) and
503 $M_{ij} = 0$ otherwise. This is broadcast across the batch and head dimensions.
504 – The softmax is applied along the last dimension (i.e., row-wise for each $T \times T$ matrix).
505 – The resulting **AttentionWeights** matrix has dimension $\mathbb{R}^{B \times H \times T \times T}$.
506 • **Step 3: Compute Head Outputs:** The attention weights are used to compute a weighted
507 sum of the value vectors.

$$\text{Heads}_{\text{out}} = \text{AttentionWeights} \cdot V$$

508 The dimensions are: $(\mathbb{R}^{B \times H \times T \times T}) \cdot (\mathbb{R}^{B \times H \times T \times d_v}) \rightarrow \mathbb{R}^{B \times H \times T \times d_v}$.

509 MHA Sub-block 3: Concatenation of Heads

510 The outputs from all H heads are combined back into a single tensor.

- 511 • **Input:** $\text{Heads}_{\text{out}} \in \mathbb{R}^{B \times H \times T \times d_v}$.
512 • **Computation:** The operation is a transpose followed by a reshape, which is equivalent to
513 concatenation.
514 1. $\text{Heads}_{\text{out}} \xrightarrow{\text{transpose}} \mathbb{R}^{B \times T \times H \times d_v}$
515 2. $\xrightarrow{\text{reshape}} \mathbb{R}^{B \times T \times (H \cdot d_v)}$
516 • **Resulting Matrix:** Since $H \cdot d_v = d_{\text{model}}$, the final concatenated matrix is:

$$\text{Concatenated} \in \mathbb{R}^{B \times T \times d_{\text{model}}}$$

517 MHA Sub-block 4: Final Linear Projection

518 The concatenated output is passed through a final linear layer to produce the MHA block's final
519 output.

- 520 • **Input:** $\text{Concatenated} \in \mathbb{R}^{B \times T \times d_{\text{model}}}$.
521 • **Weight Matrix:**
522 – $W^O \in \mathbb{R}^{d_{\text{model}} \times d_{\text{model}}}$ (rows \times columns: $d_{\text{model}} \times d_{\text{model}}$)
523 • **Computation:**
524
$$\text{MHA}_{\text{out}} = \text{Concatenated} \cdot W^O$$

525 • **Final MHA Output:** The final output of the MHA block has dimension $\mathbb{R}^{B \times T \times d_{\text{model}}}$,
526 matching the dimension of the block's input. This output is then passed to the residual
527 connection and the subsequent FFN sub-layer.

527 Addendum: A Geometric Analogy for Implicit Non-Linearity

528 The Core Idea

529 The core idea this example illustrates is that when a set of points is confined to a lower-dimensional,
530 curved surface (a manifold), a simple linear operation in the higher-dimensional “ambient” space
531 does not act linearly on the surface itself. To an observer living on the surface, the transformation
532 appears non-linear and complex.

533 This example provides an intuitive analogy for the computations within a Transformer block. The 3D
534 ambient space corresponds to the Transformer's high-dimensional embedding space ($\mathbb{R}^{d_{\text{model}}}$). The 2D
535 spherical surface represents the curved, lower-dimensional manifold onto which Layer Normalization
536 projects token vectors. Finally, the simple 3D linear transformation is analogous to the linear weight
537 matrices applied by the MHA and FFN layers.

538 A 3D Example: Shearing a Sphere

539 Let's set up our example.

540 The Setup: A 2D Surface in 3D Space

541 Imagine our set of points are not just any points in 3D space, but are all constrained to lie on the
542 surface of a sphere of radius r centered at the origin.

- 543 • **The Manifold:** This spherical surface is our 2D manifold. Although any point p on it has
544 three coordinates (x, y, z) , they are not independent.
- 545 • **The Constraint:** The coordinates must satisfy the equation:

$$x^2 + y^2 + z^2 = r^2$$

546 This constraint reduces the degrees of freedom from 3 to 2. This is analogous to how
547 LayerNorm constrains the `d_model` dimensions of a token vector.

548 The Transformation: A Linear Operation in 3D

549 Now, let's define a simple **linear transformation** that we will apply to every point in our 3D space.
550 We'll use a "shear" transformation, which shifts points in one direction by an amount proportional to
551 their coordinate in another direction.

552 Let's define a shear along the x -axis that is proportional to a point's z -coordinate. The transformation
553 matrix W and its effect on a point $p = (x, y, z)$ is:

$$W = \begin{pmatrix} 1 & 0 & s \\ 0 & 1 & 0 \\ 0 & 0 & 1 \end{pmatrix}$$

$$p' = Wp = \begin{pmatrix} 1 & 0 & s \\ 0 & 1 & 0 \\ 0 & 0 & 1 \end{pmatrix} \begin{pmatrix} x \\ y \\ z \end{pmatrix} = \begin{pmatrix} x + sz \\ y \\ z \end{pmatrix}$$

554 This is a perfectly linear transformation. It simply adds s times the z -coordinate to the x -coordinate.
555 If we apply this to a cube, it would tilt into a slanted parallelepiped.

556 The Result: A Non-Linear Transformation on the 2D Surface

557 What happens when we apply this linear shear to every point on our sphere? The sphere is distorted
558 into a new shape—a "sheared sphere," which is a type of ellipsoid.

559 The crucial question is: how can we describe the transformation that happened *on the surface*? A
560 great way to see the non-linearity is to look at the "shadow" the sphere casts on the xy -plane before
561 and after the transformation.

- 562 • **Before Transformation:** The projection of the original sphere onto the xy -plane is a filled
563 disk of radius r . A point in this disk is (x, y) .
- 564 • **After Transformation:** A point (x, y, z) on the original sphere is moved to $(x + sz, y, z)$.
565 The projection of this new point onto the xy -plane is $(x', y') = (x + sz, y)$.

566 Now, let's define the transformation purely in terms of the 2D coordinates (x, y) of the projection.
567 To do this, we must substitute for z . From our sphere's constraint equation, we know:

$$z = \pm \sqrt{r^2 - x^2 - y^2}$$

568 Substituting this into our transformation for the projected points gives us:

$$(x, y) \rightarrow (x', y') = \left(x \pm s\sqrt{r^2 - x^2 - y^2}, y \right)$$

569 This is the rule that maps a point from the original circular shadow to its corresponding point in the
570 new, distorted shadow.

571 **This transformation is clearly non-linear.** It involves square roots and squared terms. A simple,
572 clean, linear shear in 3D has induced a complex, non-linear "warping" in the 2D projection space.

Conclusion and Connection to Transformers

This example provides an intuitive analogy for what happens inside a Transformer block.

1. **LayerNorm** takes the token vectors, which could be anywhere in d_{model} -dimensional space, and projects them onto a specific, curved $(d_{\text{model}}-2)$ -dimensional manifold (analogous to our sphere).
2. The **MHA and FFN layers** then apply linear transformations (weight matrices, analogous to our shear matrix W) to these constrained vectors.
3. From the “extrinsic” view of the full d_{model} space, this operation is linear. But from the “intrinsic” view of the lower-dimensional manifold where the token representations actually live, the transformation is powerfully **non-linear**.

This interplay—projecting onto a curved surface and then applying a linear map—is a fundamental source of the Transformer’s computational expressiveness. It’s a more subtle and geometric form of non-linearity than that provided by explicit activation functions like GELU, and it happens at every sub-layer of the network.

Addendum: The Mechanics of Head Specialization

The section *Attention as Relational Computation in Specialized Bases* introduces the concept of Multi-Head Attention (MHA) as a form of relational computation in specialized bases, where each head learns to focus on a different kind of inter-token relationship. This addendum provides a more detailed, mechanistic explanation of how this specialization is achieved in a modern Transformer implementation, connecting the high-level theory to the low-level matrix operations.

Two Views of Head Projection: Conceptual vs. Implemented

There are two equivalent ways to understand the projection of input vectors into the Query, Key, and Value spaces for each head.

- **The Conceptual View:** As originally proposed, one can think of each head h as having its own distinct, smaller weight matrices: W_h^Q, W_h^K, W_h^V . The input for each head is computed via a separate projection. This view correctly emphasizes that each head has its own set of learnable parameters.
- **The Modern Implementation View:** For computational efficiency, modern implementations (as detailed in the *Addendum: DecodeBlock Pseudocode*) use a single, large weight matrix for each projection type (e.g., $W^Q \in \mathbb{R}^{d_{\text{model}} \times d_{\text{model}}}$). The input tensor X is multiplied by this large matrix once, and the resulting tensor is then reshaped to create the H heads.

These two views are computationally equivalent. The modern implementation is simply a way to perform all the per-head projections in a single, highly optimized matrix multiplication. The large weight matrix can be understood as the concatenation of the smaller, per-head weight matrices. The mechanism that makes this possible is “slicing.”

The Slicing Mechanism and its Dimensions

The `reshape` operation that follows the main projection effectively partitions the large weight matrix into dedicated “slices” for each head. Let’s analyze the Query projection as an example.

1. The MHA block takes the input tensor $X \in \mathbb{R}^{B \times T \times d_{\text{model}}}$ and multiplies it by the full query weight matrix $W^Q \in \mathbb{R}^{d_{\text{model}} \times d_{\text{model}}}$.
2. This produces a tensor of pre-queries, $Q' = X \cdot W^Q$, which has the shape $\mathbb{R}^{B \times T \times d_{\text{model}}}$.
3. The `reshape` operation then splits the final dimension, d_{model} , into H heads, each of dimension d_k (where $d_{\text{model}} = H \times d_k$). This means the first d_k columns of the Q' tensor are allocated to Head 1, the next d_k columns to Head 2, and so on.

617 The crucial insight is that the first d_k columns of the output tensor Q' are computed by multiplying
 618 the input X with only the first d_k columns of the weight matrix W^Q . This sub-matrix of W^Q is the
 619 effective weight matrix for Head 1.

620 We can therefore define the per-head weight **slices** as follows:

- 621 • **Query Slice for Head h :** $W_{\text{slice}_h}^Q \in \mathbb{R}^{d_{\text{model}} \times d_k}$
- 622 • **Key Slice for Head h :** $W_{\text{slice}_h}^K \in \mathbb{R}^{d_{\text{model}} \times d_k}$
- 623 • **Value Slice for Head h :** $W_{\text{slice}_h}^V \in \mathbb{R}^{d_{\text{model}} \times d_v}$

624 For each head h , the set of slices $(W_{\text{slice}_h}^Q, W_{\text{slice}_h}^K, W_{\text{slice}_h}^V)$ constitutes its unique, learnable parameters.
 625 These are the “specialized linear projections” mentioned in Section 4.3.

626 Specialization through Independent Gradient Flow

627 The specialization of heads is a direct and mechanical consequence of how gradients are calculated
 628 during backpropagation.

- 629 1. **Forward Pass:** During the forward pass, the computation for each head is independent. The
 630 attention scores and output for Head 1 depend only on the projections generated by its own
 631 weight slices, $(W_{\text{slice}_1}^Q, W_{\text{slice}_1}^K, W_{\text{slice}_1}^V)$.
- 632 2. **Backward Pass (Backpropagation):** When the model’s final loss is calculated, the gradient
 633 is propagated back through the network. The chain rule of calculus ensures that the gradient
 634 signal for a given head’s computation flows back *only* to the parameters that were involved
 635 in that computation.
- 636 3. **Independent Optimization:** This means the loss gradient attributable to Head 1’s per-
 637 formance will only update the weights within its dedicated slices. The weights in $W_{\text{slice}_2}^Q$
 638 will not be affected by the gradient from Head 1. This creates H independent optimization
 639 pathways.
- 640 4. **Emergence of Specialists:** Over millions of training steps, this process naturally leads to
 641 specialization. If the model consistently makes errors on syntactic tasks that Head 1 “should
 642 have” captured, the gradients will preferentially update the weights in Head 1’s slices to
 643 make it a better “syntax detector.” Simultaneously, if Head 2 fails on semantic tasks, its
 644 slices will be optimized to create a better “semantic basis.”

645 Conclusion: Specialization as a Mechanical Consequence

646 Head specialization is not a magical emergent property but a direct result of the architecture’s
 647 design. By parameterizing each head with its own unique and independently optimized set of weights
 648 (implemented as slices of a larger matrix), the architecture creates parallel computational workspaces.
 649 The process of backpropagation then naturally tunes each workspace to become a specialist at
 650 minimizing the types of errors it is most responsible for, resulting in a team of experts that collectively
 651 analyze the input sequence from multiple, complementary perspectives.

652 Addendum: Summary of Theoretical Contributions

653 This paper presents a new, unified framework for understanding the Transformer architecture. Our
 654 contributions are not empirical but theoretical, derived from a first-principles analysis of the geometry
 655 and information theory of the model’s components. The following tables summarize our main results:
 656 **Table 1: Core Theoretical Results**, **Table 2: Answers to Open Questions**, and **Table 3: New**
 657 **Avenues for Research**.

658 The rest of this page may be blank, see next page.

Table 1: Core Theoretical Results of the Framework.

Theoretical Result	Significance
Layer Normalization (LN) is a geometric projection operator forcing representations onto a $(d_{\text{model}} - 2)$ -dimensional manifold.	Reframes LN from an ad-hoc trick to a fundamental architectural principle.
Linear transformations on the post-LN curved manifold become powerful non-linear operators.	Reveals a pervasive and previously unappreciated source of the Transformer’s expressive power that exists in every sub-layer.
The attention mechanism’s dot product is a direct measure of the geodesic distance on the post-LN hypersphere.	Demonstrates a perfect, principled alignment between the geometry of the representational space and the core computational primitive.
The fixed LN manifolds act as geometric stabilizers against representational drift in practical, sliding-window transformers.	Provides the first concrete, mechanistic explanation for the architecture’s remarkable long-context coherence.
The Transformer block operates on a symmetrize-and-structure principle (isotropic projection followed by anisotropic processing).	Replaces the "bag of tricks" view of the architecture with a coherent, repeating design principle.
Multi-Head Attention’s (MHA) axis-aligned partitioning incentivizes the network to encode meaning in vector axis indices.	Proposes a new, structured model for how representations are organized, suggesting a deeper level of learned order in the latent space.
The geometric mismatch between the $(d_{\text{model}} - 2)$ -dimensional data manifold and the d_{model} -structured processing grid of MHA acts as a powerful implicit regularizer.	Provides a novel, geometric explanation for the model’s robustness and generalization.
The Feed-Forward Network (FFN) acts as a manifold-based information filter performing a complexity-reducing deformation.	Provides a clear, information-theoretic purpose for the FFN’s specific two-layer structure.
The symmetrize-and-structure cycle is isomorphic to a Renormalization Group (RG) flow from fundamental physics.	Dissolves the "unreasonable effectiveness" of the Transformer by showing it has convergently discovered a universal strategy for analyzing multi-scale systems.

Table 2: Answers to Open Questions in Transformer Theory.

Open Question Addressed	Answer Provided by this Work	Significance of the Answer
What is the true role of Layer Normalization beyond statistical stabilization?	It is a geometric projection operator forcing representations onto a $(d_{\text{model}} - 2)$ -dimensional manifold.	Reframes LN from an ad-hoc trick to a fundamental architectural principle.
What are the origins of the Transformer’s expressive power beyond explicit activations?	Linear transformations on the post-LN curved manifold become powerful non-linear operators.	Reveals a pervasive, unappreciated source of non-linearity in every sub-layer.
What is the nature of the synergy between Layer Normalization and the attention mechanism?	The attention dot product is a direct measure of geodesic distance on the post-LN hypersphere.	Demonstrates a perfect alignment between the representational geometry and the computational primitive.
How do Transformers maintain coherence over long contexts?	The fixed LN manifolds act as geometric stabilizers against representational drift in sliding-window transformers.	Provides the first concrete, mechanistic explanation for the architecture’s remarkable stability.
Does the Transformer block possess a coherent design philosophy?	It operates on a repeating symmetrize-and-structure principle.	Replaces the "bag of tricks" view with a unified design principle.
What is the mechanism behind the effectiveness of Multi-Head Attention?	Its axis-aligned partitioning creates a strong incentive for the network to encode meaning in the indices of its vector axes.	Proposes a new, structured model for how representations are organized in the latent space.
What is the source of the Transformer’s emergent robustness and generalization?	The geometric mismatch between the $(d_{\text{model}} - 2)$ -dimensional data manifold and the d_{model} -structured processing grid of MHA acts as a powerful implicit regularizer.	Provides a novel, geometric explanation for the model’s ability to generalize.
What is the functional purpose of the FFN’s specific two-layer structure?	The FFN acts as a manifold-based information filter that performs a complexity-reducing deformation.	Provides a clear, information-theoretic purpose for the FFN’s design.
Why is the Transformer architecture "unreasonably effective"?	Its core symmetrize-and-structure cycle is isomorphic to a Renormalization Group (RG) flow from fundamental physics.	Dissolves the mystery of the architecture’s success by showing it has convergently discovered a universal strategy for analyzing complex, multi-scale systems.

Table 3: New Research Avenues Opened by the Framework.

New Research Avenue	Enabling Theoretical Result	Importance of the New Avenue
Geometric Design of Normalization Layers	LN is a projection operator onto a fixed, $(d_{\text{model}} - 2)$ -dimensional manifold.	Provides a new, geometric language for analyzing and designing normalization layers, moving beyond purely statistical justifications.
Sources of Implicit Non-Linearity	Simple linear layers become powerful non-linear operators when their inputs are constrained to the curved manifold created by LN.	Suggests architectural design can be as important as explicit activation functions for creating non-linear systems.
Co-design of Latent Space Metrics and Operators	The attention dot product is a direct measure of geodesic distance on the post-LN hypersphere.	Suggests future architectures can be designed by explicitly aligning the geometry of the representation with the algebra of the computation.
Internal Structure of Learned Representations	The axis-aligned partitioning of MHA creates a strong incentive for the network to encode meaning in the indices of its vector axes.	Provides a new, testable model for how Transformers organize information, potentially leading to more interpretable models.
Implicit Geometric Regularization	The mismatch between the $(d_{\text{model}} - 2)$ -dimensional data manifold and the d_{model} -structured processing grid of MHA acts as a powerful regularizer.	Identifies a novel mechanism for improving generalization, suggesting "incommensurability" could be a new design principle for robust architectures.
Learnable and Dynamic Architectural Components	The hard partitioning of MHA can be generalized to continuous, learnable "soft head" functions.	Transforms a fixed architectural choice into a dynamic, learnable parameter, paving the way for adaptive models.
Connection between Deep Learning and Physics	The symmetrize-and-structure cycle is isomorphic to a Renormalization Group (RG) flow.	Provides a powerful new language and conceptual toolkit (fixed points, universality classes) for analyzing hierarchical learning.
Biologically Plausible Normalization	The proposal for a trainable, anisotropic LayerNorm—a learnable, iterative, multi-headed block.	Provides a concrete path toward designing components that are both architecturally consistent and more aligned with biological principles.